

YIJIE (EJ) ZHOU

yz926@cam.ac.uk ♦ ej-zhou.github.io ♦ [GitHub](#) ♦ [Google Scholar](#)

Research Interests: NLP - Interpretability, Generalisation, Multilinguality, Bias & Calibration

EDUCATION

University of Cambridge , PhD in Computation, Cognition and Language <i>Language Technology Lab</i> , Supervised by Anna Korhonen	Cambridge, UK Oct 2024 – Present
University of Oxford , Visiting Student <i>St Hilda's College</i> , Linguistics and Computer Science	Oxford, UK Oct 2023 – Jun 2024
Cornell University , Exchange Student & Research Assistant <i>College of Engineering & CommCollabTech Lab</i> , GPA: 4.082/4.3	Ithaca, NY, USA Jan 2023 - May 2023
Zhejiang University , Bachelor (<i>Hons</i>) of Computer Science <i>Chu Kochen Honors College</i> , Overall GPA: 3.99/4, Major GPA: 4.0/4.0, Top 2%	Hangzhou, China Sept 2020 - Jun 2024

EXPERIENCE

Alibaba Qwen , Research Intern <i>Qwen Multilingual Team</i> , Advisor: Baosong Yang	Hangzhou, China Jul 2026 – Present
Yale University , Research Assistant <i>Yale NLP Lab</i> , Advisor: Arman Cohan	New Haven, CT, USA Apr 2023 - Oct 2023
Westlake University , Research Intern <i>WestlakeNLP Lab</i> , Advisor: Yue Zhang & Yulong Chen	Hangzhou, China Sept 2022 - Feb 2023

SELECTED PUBLICATION

- [1] **Ej Zhou**, Lucas Resck, Zheng Hui, Anna Korhonen: *LLM Evaluators are Biased towards Languages*, in **COLM 2026** & in ICLR 2026 ICBINB.
- [2] **Ej Zhou***, Songbo Hu*, Yinhong Liu*, Evgeniia Razumovskaia, Xiaobin Wang, Alexander Fraser, Ivan Vulić, and Anna Korhonen: *Dial HEALTHDIAL for Advice: A Multilingual and Multi-Parallel Spoken Dialogue Dataset for Knowledge-Grounded Information Seeking* [\[paper\]](#) [\[code\]](#) [\[video\]](#), in **ACL 2026** Findings.
- [3] Jiaqi Weng, Han Zheng, Hanyu Zhang, **Ej Zhou**, Qinqin He, Jialing Tao, Hui Xue, Zhixuan Chu, Xiting Wang: *Safe-SAIL: Towards a Fine-grained Safety Landscape of Large Language Models via Sparse Autoencoder Interpretation Framework* [\[paper\]](#), in **ACL 2026** Findings.
- [4] Pedro Ortiz Suarez **et al.** (incl. **Ej Zhou**): *CommonLID: Re-evaluating State-of-the-Art Language Identification Performance on Web Data* [\[paper\]](#) [\[blog\]](#), contributing author, in **ACL 2026**.
- [5] **Ej Zhou**, Weiming Lu: *Bias Beyond English: Evaluating Social Bias and Debiasing Methods in a Low-Resource Setting* [\[paper\]](#), in **NLPCC 2025** & ACL 2025 GeBNLP.
- [6] Rohan Phanse, **Yijie Zhou**, Kejian Shi, Wencai Zhang, Yixin Liu, Yilun Zhao, Arman Cohan: *MSRS: Evaluating Multi-Source Retrieval-Augmented Generation* [\[paper\]](#) [\[code\]](#), in **COLM 2025**.
- [7] Yulong Chen, Huajian Zhang, **Ej Zhou**, Xuefeng Bai, Yueguan Wang, Ming Zhong, Jianhao Yan, Yafu Li, Judy Li, Michael Zhu and Yue Zhang: *Revisiting Cross-Lingual Summarization: A Corpus-based Study and A New Benchmark with Improved Annotation* [\[paper\]](#) [\[code\]](#), in **ACL 2023**.

Workshops & under review:

- [8] **Ej Zhou**, Suchir Salhan, Catherine Arnett, Anna Korhonen: *Cross-Lingual Alignment Without Joint Training: Do Monolingual Language Models Converge on Universal Representations?*, in ICLR 2026 UCRL & in submission to EMNLP 2026.

[9] **Ej Zhou**, Caiqi Zhang, Tiancheng Hu, Chengzu Li, Nigel Collier, Ivan Vulić, Anna Korhonen: *Beyond the Final Layer: Intermediate Representations for Better Multilingual Calibration in Large Language Models* [paper], **spotlight talk** in COLM 2025 MELT & in NeurIPS 2025 MechInterp, in submission to EMNLP 2026.

[10] **Ej Zhou***, Zijie Zheng*: *What If Chinese Were Latinized? A Counterfactual Study of Script, Tokenization, and Language Modeling*, in ICML 2026 Culture x AI & in submission to EMNLP 2026.

[11] Zheng Hui, Sanhanat Sivapiromrat, **Ej Zhou**, Arnau Marin-Llobet, Nigel Collier: *Rethinking Intent Concealment: When Jailbreak Wrappers Backfire on Aligned LLMs*, in submission to EMNLP 2026.

[12] Giulia Occhini, Kumiko Tanaka-Ishii, Anna Barford, Refael Tikochinski, Songbo Hu, Roi Reichart, **Yijie Zhou**, Hannah Clause, Ulla Petti, Ivan Vulić, Ramit Debnath, Anna Korhonen: *Artificial Intelligence is Creating a New Global Linguistic Hierarchy* [paper] [article], in submission to Nature Communications.

* denotes equal contribution.

SERVICE & EVENTS

Teaching - Supervisor & Lecturer, Li 18 Computational Linguistics 2025-2026, Cambridge

Conferences — ICML 2026, EACL 2026, EMNLP 2025, ACL 2025, NLPCC 2025, ACL 2024, EACL 2024, EMNLP 2023, ACL 2023

Volunteering - EAG London 2026, EACL 2024

Committee Member - Cambridge World Cinema Society, Cambridge Language and Culture Society

Events — MLSS^{R&S} 2026, Cambridge Language Sciences Annual Symposium (Posters*2), SparkLab S1 (Mentor), OxGen 2023 Summit & OxGen AI Hackathon 2023 (Special Award), Polyglot Conference 2023, Language Event 2023

SELECTED AWARDS

- | | |
|--|-----------|
| • Cambridge Trust Scholarship (Fully Funded) | 2024-2028 |
| • Excellence Honors Scholarship, Zhejiang University (<i>Top 1% in CKC Honors College</i>) | 2021 |
| • Hengyi Scholarship (<i>5/~24,000 in Zhejiang University</i>) | 2021 |

LANGUAGES

- I do multilingual research because I love languages. I speak to varying degrees: *Mandarin, English, Wu, French, Russian, Japanese, German*; I'm learning *Polish* and *Portuguese*.
- I have a 800+ day streak on Duolingo.